

Feature Extraction for Deterministic Dynamical Systems

A chaos theory perspective on predictive maintenance

Michael Small, Ayham Zaitouny and Débora Corrêa



Outline

Chaos:

Why Chaos? What data?

Embedding

Tipping points

Surrogates

Further reading

Recurrence:

Embedding review

Recurrence Plots

Recurrence Quantification
analysis (RQA)

RPs to detect change points

Ordinal Partition Networks:

Recurrence Networks review

Ordinal Partition Networks

Permutation Entropy

Application to data

Part I

Chaos in predictive maintenance

Why Chaos? What data?

What is Chaos?

- The basic task in predictive maintenance is to preempt change in the system behaviour
- Achieved via either stochastic (i.e. Bayesian or frequentist) or deterministic (machine learning or chaotic) models
- Build a model a check when the model stops working \Rightarrow system is changing

Example (sensitivity to initial conditions)

<https://www.youtube.com/watch?v=n-mpifTiPV4>

Example (chaos in a mechanical double pendulum)

<https://www.youtube.com/watch?v=U39RMUzCjiU>

Bernoulli map – the simplest example of Chaos

Simplest example of chaos (and specifically sensitive dependence on initial conditions):

$$x_{n+1} = f(x_n) = \begin{cases} 2x_n & 0 \leq x_n < 0.5 \\ 2x_n - 1 & 0.5 \leq x_n \leq 1 \end{cases}$$

Hence:

- This map is bounded by (and onto) $[0, 1]$
- $f'(x) = 2$ a.e. and so local separation increases as 2^n .
- No $x_0 \in \mathbb{R} \setminus \mathbb{Q}$ is periodic \implies chaos a.e.
- Every $x_0 \in \mathbb{Q}$ is periodic \implies periodic orbits are dense

Note, that this means that while almost every random initial condition leads to chaos, every initial condition chosen by a computer is periodic (with period 1).

Logistic growth

In continuous time,

$$x'(t) = \lambda x(t)(1 - (x(t)))$$

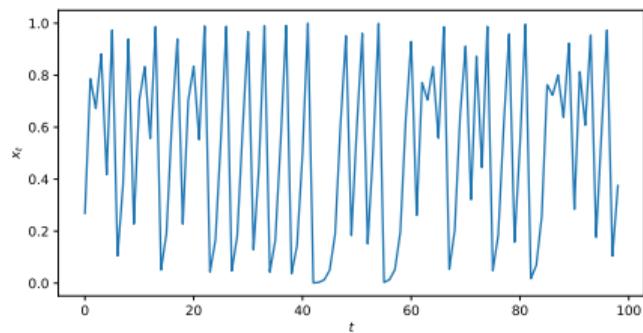
is a simple paradigmatic model of population growth with capacity. In discrete time,

$$x_{t+1} = \lambda x_t(1 - x_t)$$

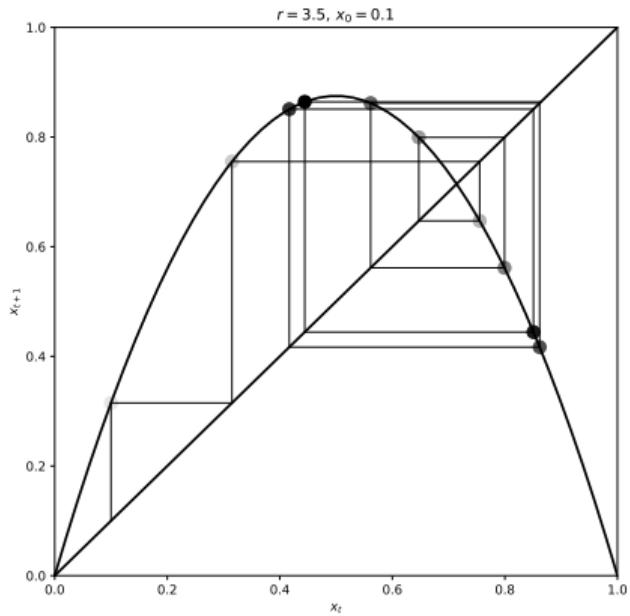
becomes chaotic and is (perhaps) one of the simplest examples of non-trivial practical chaotic dynamics. There are two ways of understanding chaotic dynamical systems like this:

- the *bifurcation diagram*; and,
- *phase space*.

Phase space

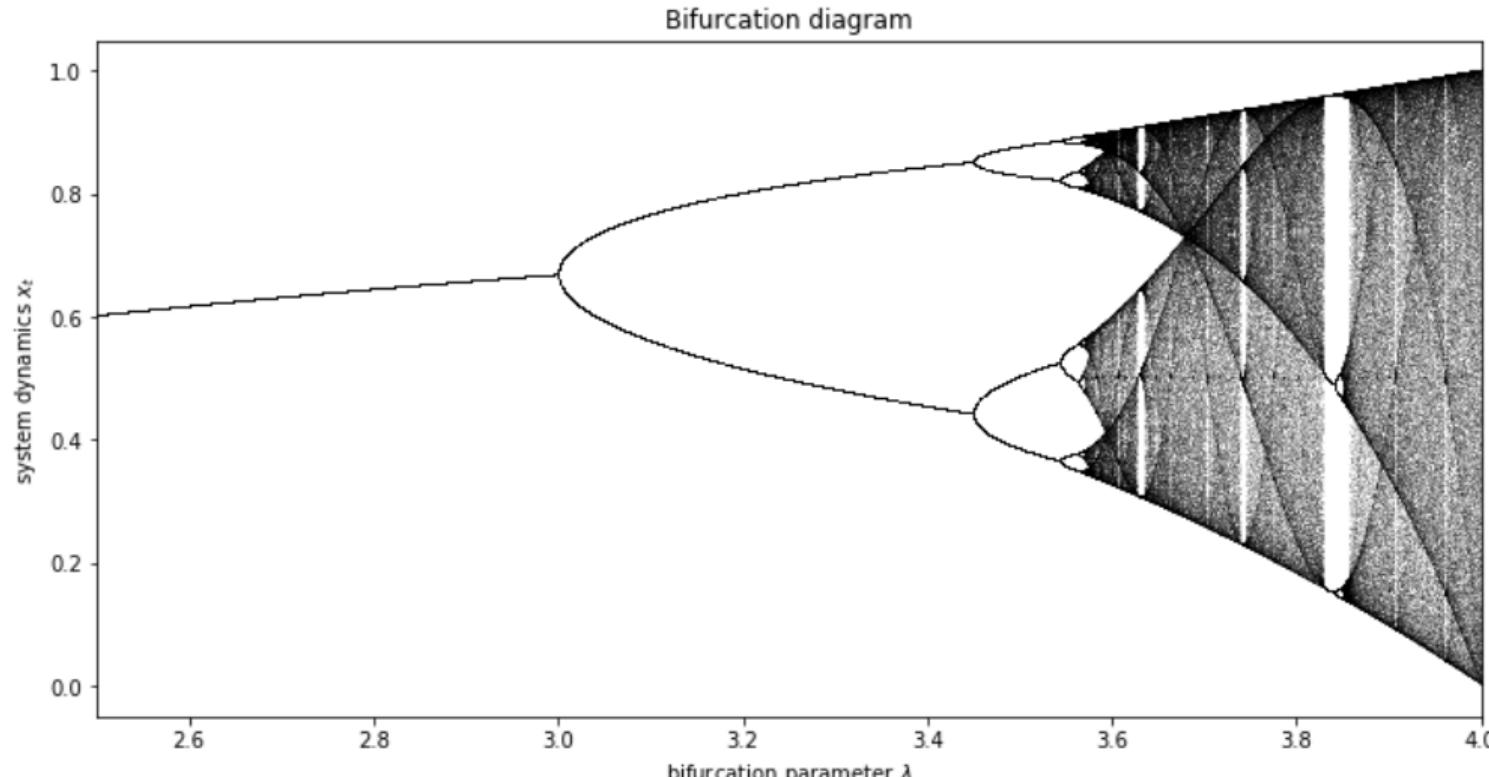


the time series



phase space

Bifurcation



the bifurcation diagram

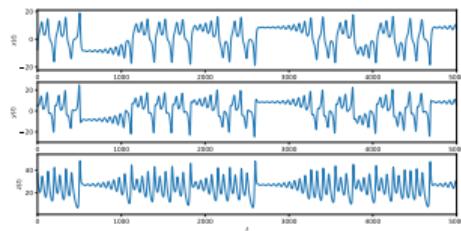
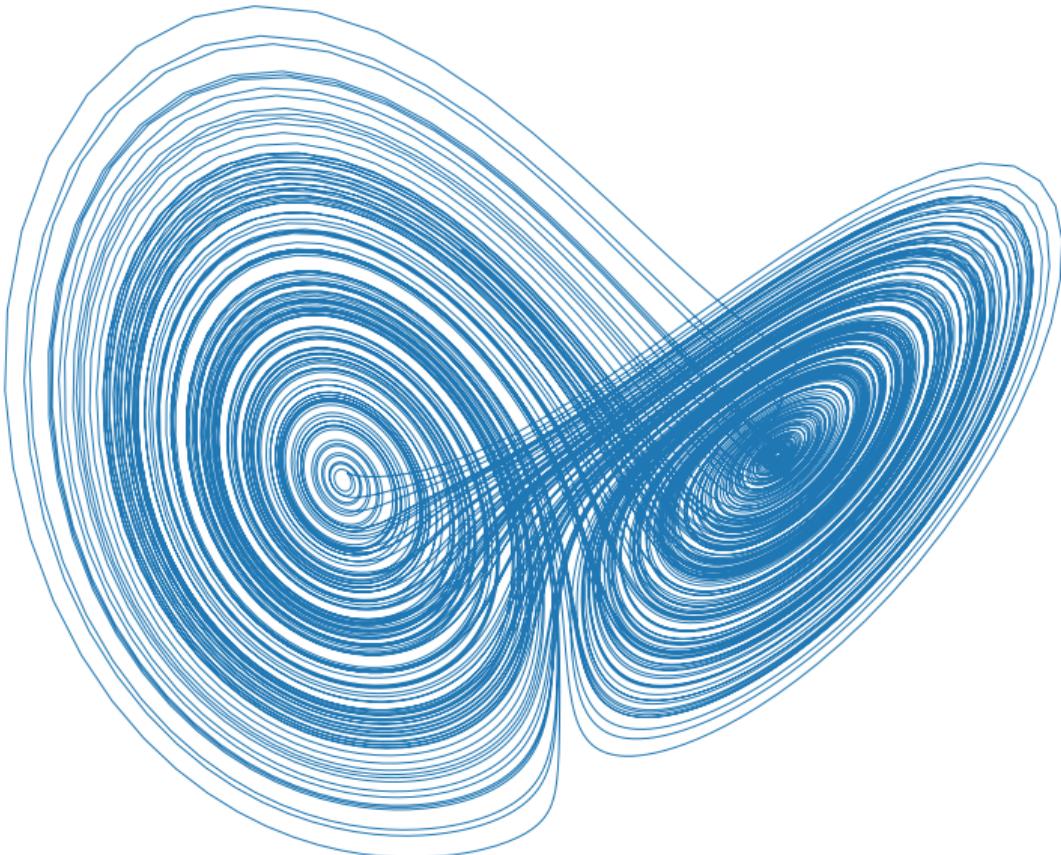
The Lorenz “butterfly”

Lorenz proposed a three dimension set of ordinary differential equations to model convection forced via a temperature gradient:

$$\begin{aligned}\frac{dx}{dt} &= -\sigma x + \sigma y \\ \frac{dy}{dt} &= -xz + rx - y \\ \frac{dz}{dt} &= xy - bz\end{aligned}$$

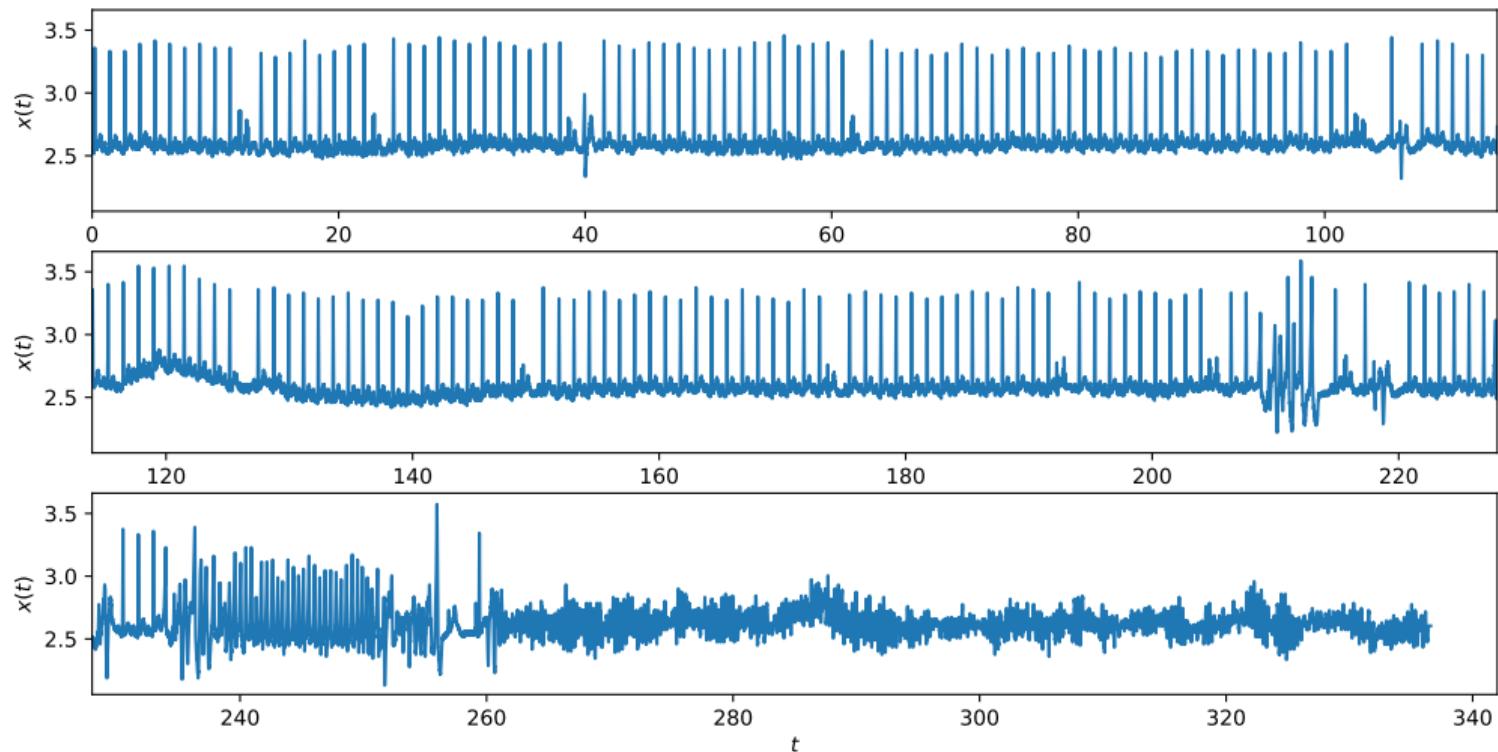
where (it turned out), for $\sigma = 10$, $r = 28$, and $b = \frac{8}{3}$, the system is bounded and aperiodic. The variable x represents intensity of convective motion, y temperature gradient, and z is the degree of nonlinearity in the temperature profile.

Lorenz in phase space

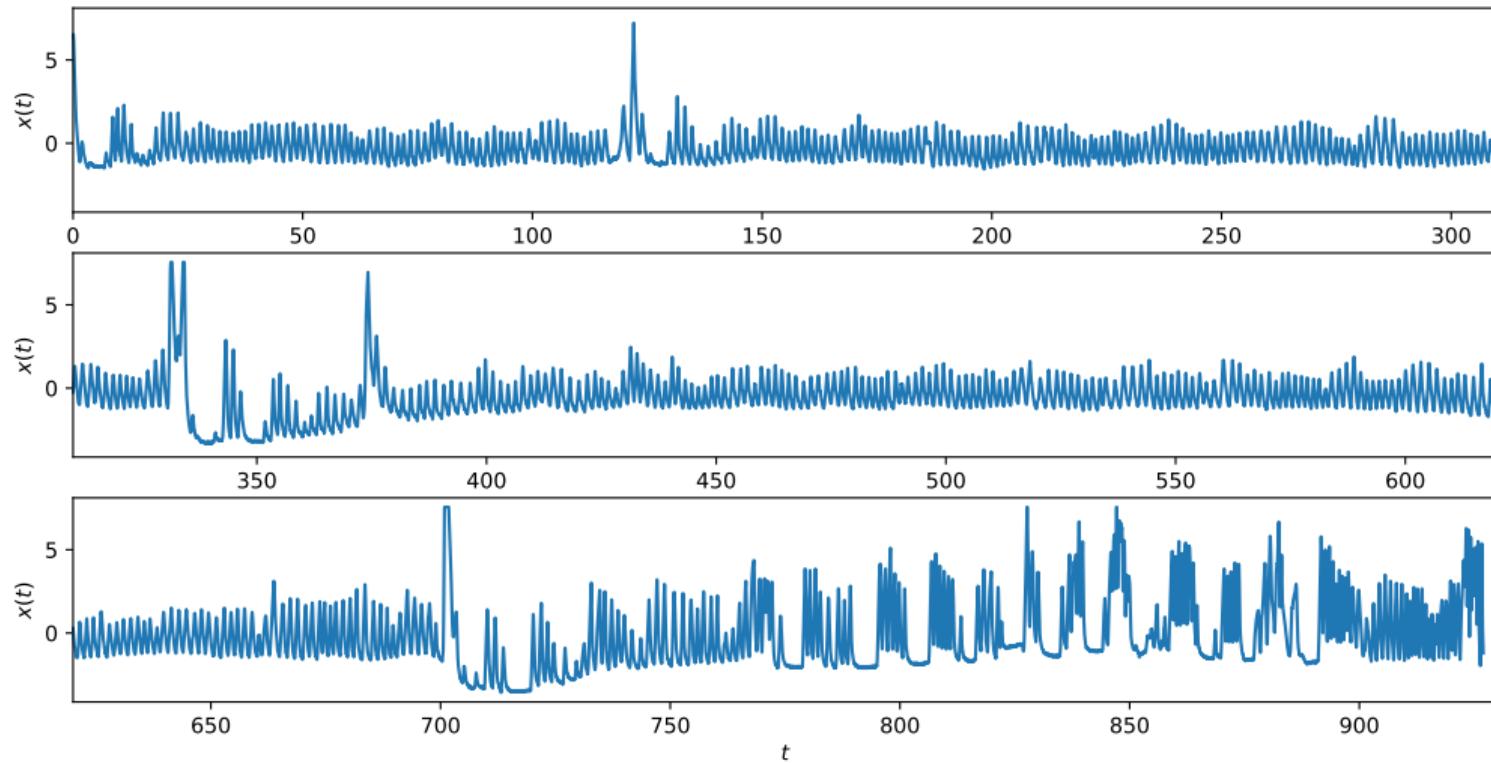


- What are the characteristics of a “good” data set?
- What sort of data are generated in industry ?

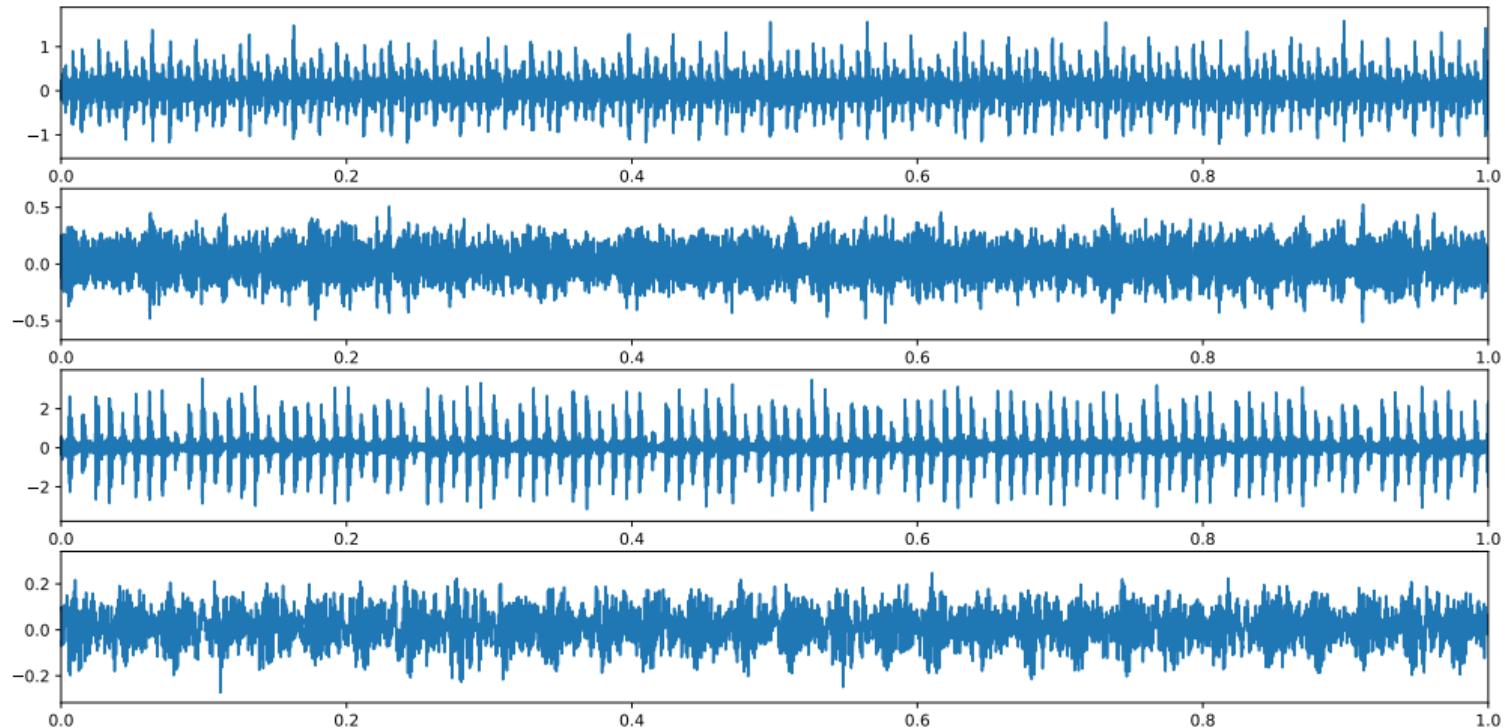
Electrocardiogram during a heart attack



Inductance Plethysmography during onset of Cheynes-Stokes respiration



Mechanical bearing failure



Open data available from csegroups.case.edu



The Complex Systems Group

Embedding

An Experiment

Let $\Phi(z_0) = \{\phi_t(z_0) | t \in T\}$ be the *trajectory* for an initial condition z_0 ¹

Let $h : \mathcal{M} \rightarrow \mathbb{R}$ be a *measurement function* and suppose that we can only observe

$$x_i = h(\phi_{i\kappa}(z_0))$$

for $i \in \mathbb{Z}^+ \cup \{0\}$. We call κ the *sampling rate* of our experiment.

What can the *time series* $\{x_i\}_{i=1}^N$ tell us about ϕ ?

¹We can, of course, think about trajectories backwards in time as well – and to do so is more typical. Moreover, we have yet to prove that $\Phi(z_0)$ is well defined, but suppose for now that it is.

Takens' Embedding Theorem (1981)

The map

$$x_i \mapsto (x_i, x_{i+1}, \dots, x_{(i+m-1)}) =: v_i$$

is an embedding of a compact manifold with dimension $d \in \mathbb{Z}^+$ ($m = 2d + 1$) if $h : \mathcal{M} \rightarrow \mathbb{R}$ is C^2 and “generic”². Moreover, evolution of v_i are diffeomorphic to the dynamics of ϕ .

Corollary (Sauer, York, and Casdagli, 1981)

Takens' embedding theorem also holds for $d \in \mathbb{R}^+$ with the condition that $m \geq 2d + 1$.

Hence, a fractal attractor $\mathcal{A} \subset \mathcal{M}$ can be *reconstructed* from a time series generated from a trajectory lying on that attractor.

²Sufficiently well coupled between the d -dimensional variables and the theorem is then true almost always.

Delay reconstruction of experimental data

Suppose that a time series, as defined above, is the output of a deterministic and stationary (autonomous) dynamical system.

Techniques to reconstruct the underlying attractor are now well established³ and widely used.

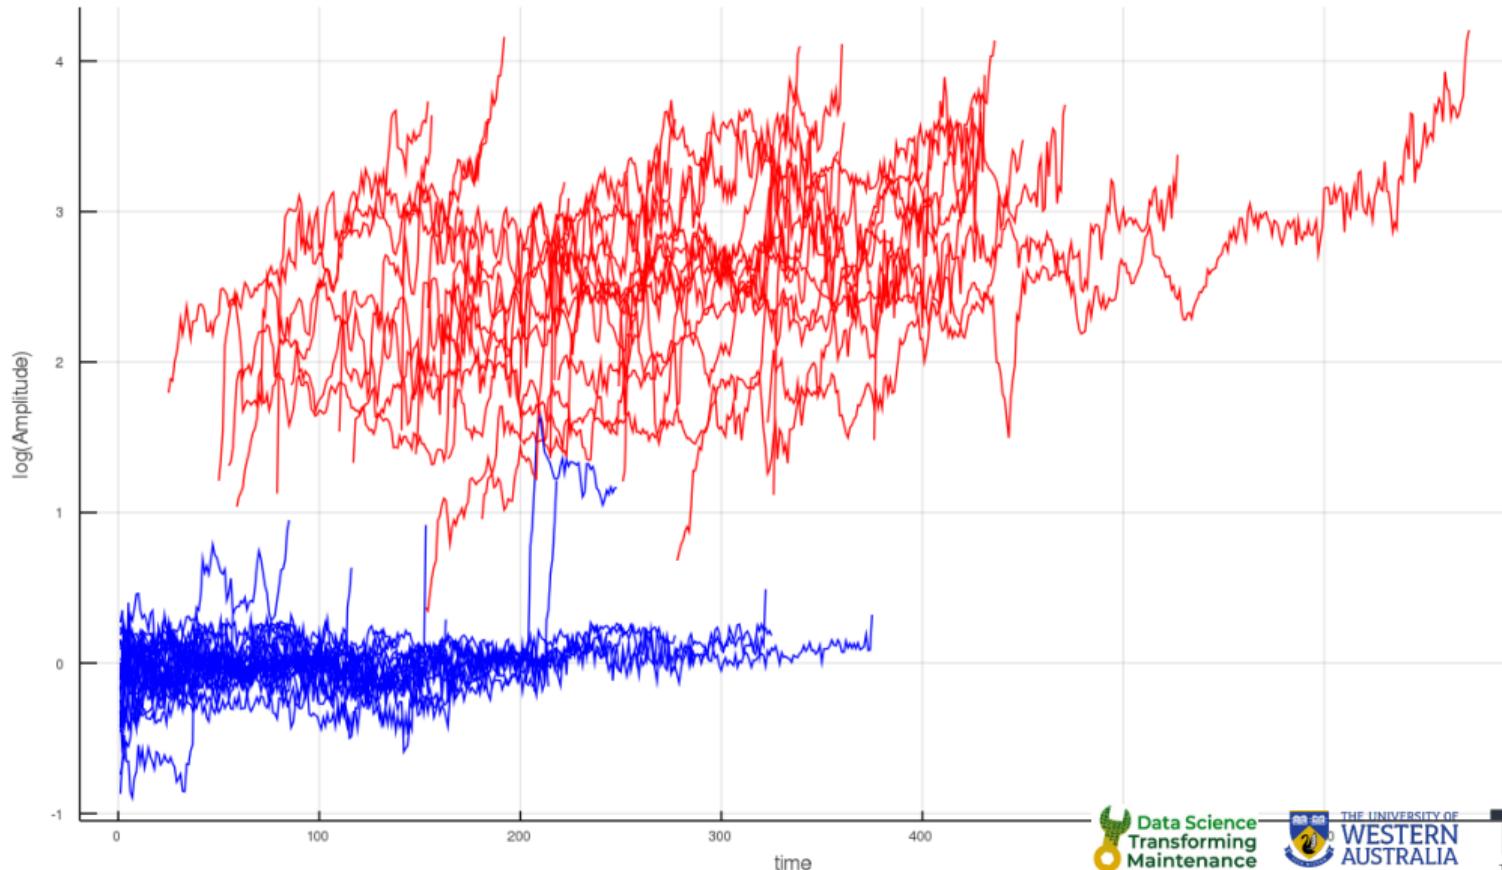
This is usually achieved by estimating two parameters: *embedding dimension m*, and *embedding lag τ* and applying the map

$$x_i \mapsto (x_i, x_{i-\tau}, \dots, x_{(i-(m-1)\tau)}) =: v_i$$

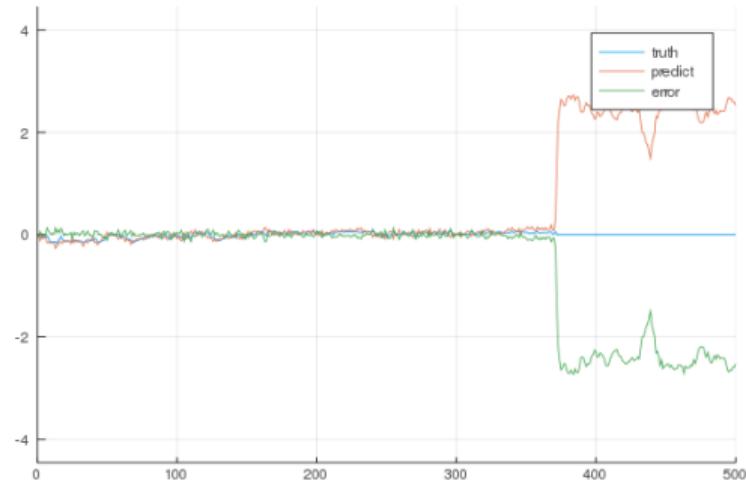
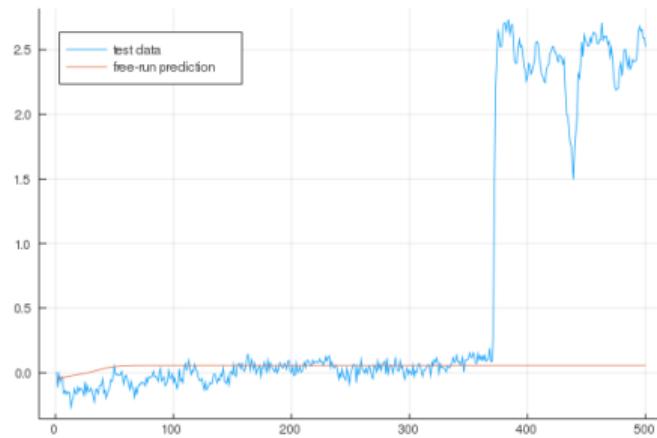
³See M. Small Applied Nonlinear Time Series Analysis, World Scientific, 2003.

Tipping points

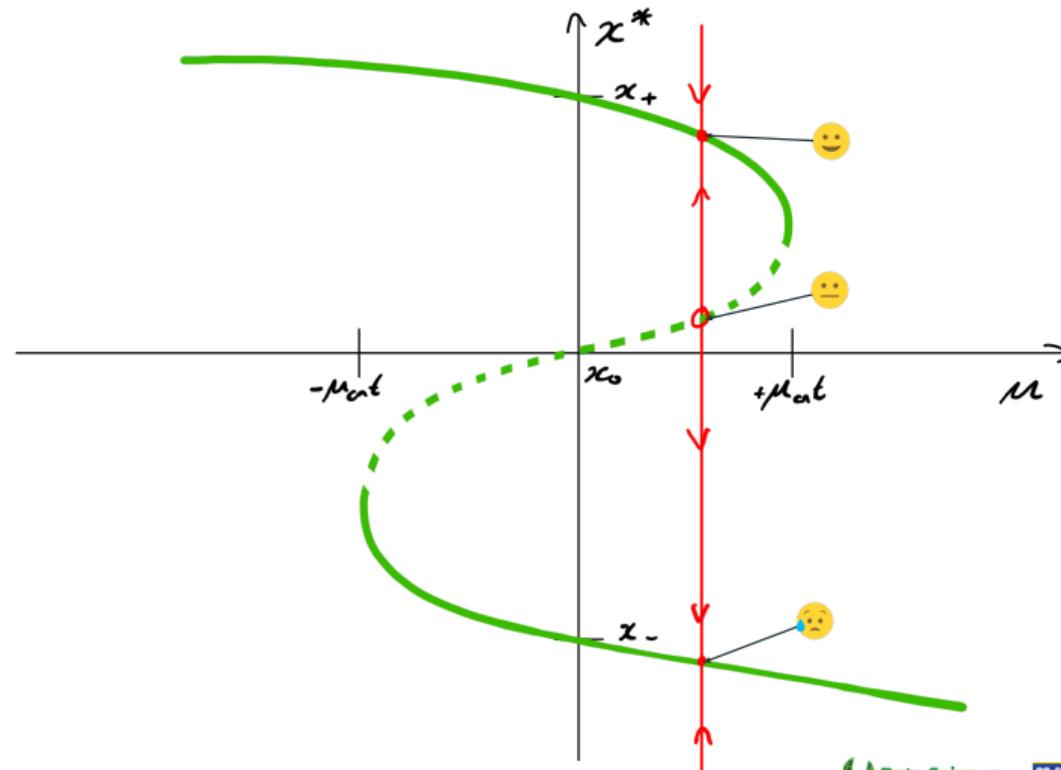
Some maintenance data: Component failure



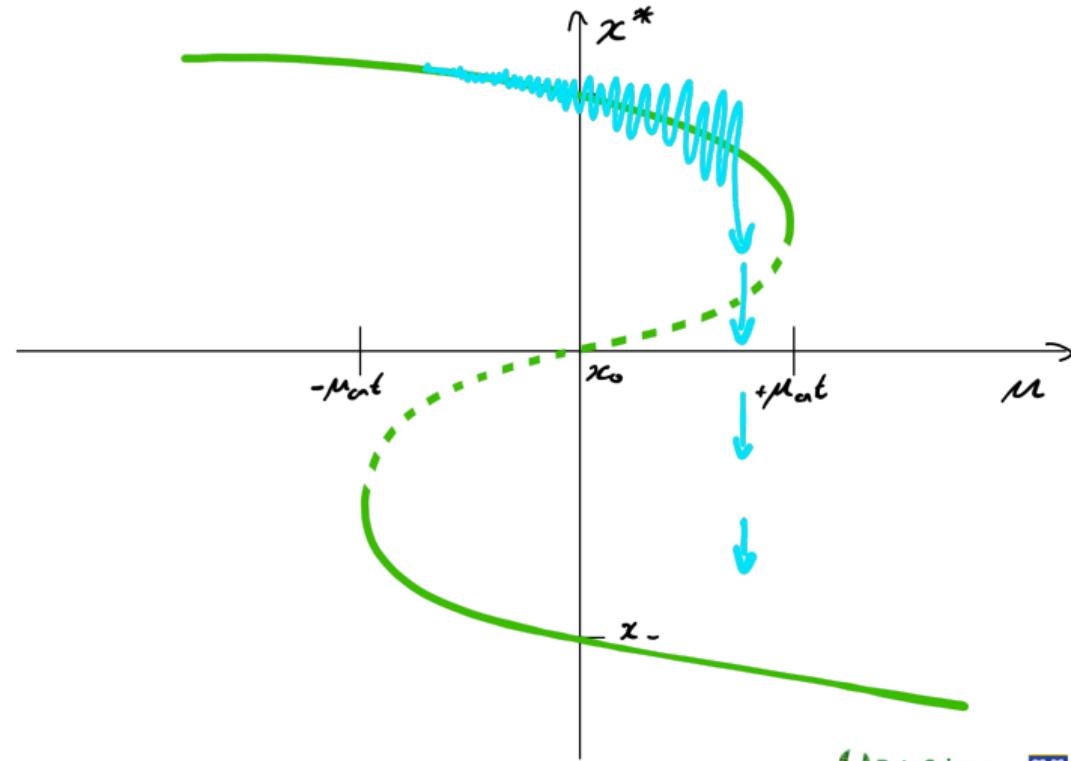
Model it



Hysteresis and a tipping point



Hysteresis and a tipping point



The tipping point model of dynamics

The signature of tipping point transitions are

- loss of stability
- increase sensitivity to noise
- increase error in predictive model (built on historical data)

Surrogates

Surrogate data

- How can we know that an estimated statistical quantity (hopefully an invariant of the underlying dynamical system) has been reliably estimated from data?
- What does the estimated correlation dimension (Lyapunov exponent, etc.) actually mean?
- What would we expect for boring (linear noise) data?
- Are nonlinear time series techniques warranted by the data? Or would linear methods suffice?
- Which class of model (linear vs. nonlinear, radial basis functions, etc.) are warranted by, and produce results consistent with, the data?

Hypothesis testing

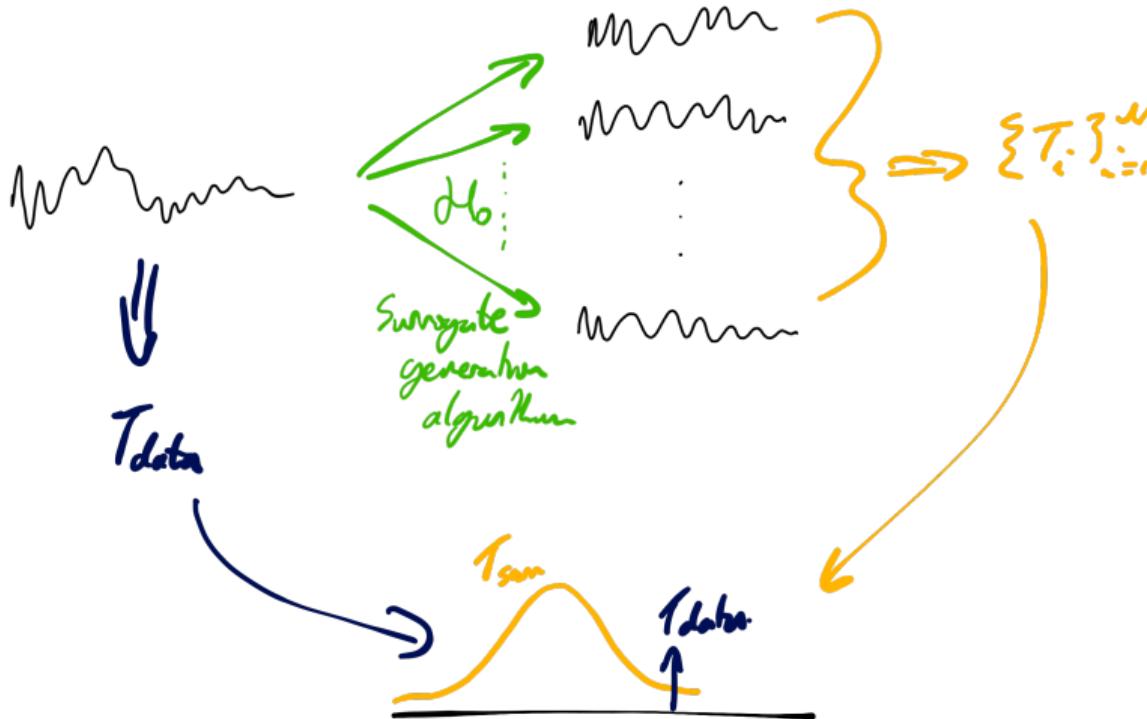
Definition

Hypothesis testing in statistics is the idea that one can compute some statistical value from observed data, compare that value to some theoretical distribution and conclude that an underlying (null) hypothesis is either *rejected* or that we *fail to reject* said hypothesis.

Definition

A *surrogate data hypothesis test* consists of two components – (i) an *algorithm* with which to generate surrogate (that is, randomised) realisations s_t from an observed time series x_t , and (ii) a *test statistic* $d(\cdot)$. One computes values of the test statistic for the data $d_* := d(\{x_t\}_t)$ and for the ensemble of (many) surrogates $d_k := d(\{s_t\}_t)$. Comparing the statistic values for the data d_* to the distribution generated from the surrogates $\{d_k\}_k$, one can then make the decision to “reject” or “fail to reject”.

The surrogate algorithm



Example (Theiler's "Algorithm 0")

Let $\{x_t\}_{t=1}^N$ be an observed time series, and then an *algorithm 0* surrogate realisation $\{s_t\}_{t=1}^N$ is obtained by resampling, with replacement: choose a bijection $\tau : \mathbb{Z}_n \mapsto \mathbb{Z}_n$ then $s_t = x_{\tau(t-1)+1}$ for $t = 1, 2, 3, \dots, N$.

Note:

- s_t is a random reordering (shuffling) of x_t
- s_t is not correlated with $s_{t-\ell} \forall \ell$.

Example (applied to linear noise)

Suppose that x_t is a realisation of an autoregressive linear noise process and that the autocorrelation with lag one $\rho(1) = E(x_t x_{t-1}) =: \rho_1 \neq 0$. Let $d(\cdot) = E(x_t x_{t-1})$, then $d_* = \rho_1$ and $E(d_k) = 0$ (the d_k follow a normal distribution with mean 0 and variance scaling with $\frac{1}{k}$). For k sufficiently large one can reject the hypothesis that x_t is independent and identically distributed noise.

Algorithm 0 tests the hypothesis that the observed data is *independent and identically distributed noise*.

Example (“Algorithm 1”: Fourier transform surrogates)

Let X_n be the discrete Fourier transform of x_t (which we denote $\mathcal{F}(\{x_t\}) = \{X_n\} = \{R_n e^{-i\pi\theta_n + \phi_n}\}$). The Fourier transform surrogate is the inverse Fourier transform of X_n with the complex phases randomised (pairwise, to ensure that the result is real) over 2π :

$$s_t = \mathcal{F}^{-1} \left(\{R_n e^{-i\pi\theta_n + \phi_n + \psi_n}\}_n \right)$$

where ψ_n are uniform random numbers on $[0, 2\pi]$.

Algorithm 1 tests the hypothesis that the observed data is *linearly filtered noise* (NOTE: this means that we should suppose that the x_t have a Gaussian distribution).

Inherent in the Fourier transform are all the assumption about linearity and periodicity. When these assumption are violated, one expects to reject the hypothesis, however there is no guarantee that something loopy might not happen. Moreover, this scheme works when the input data is (close enough to) Gaussian distributed. Otherwise, what is being tested should be trivially rejected and the following test is better.

Example (“Algorithm 2”: Amplitude adjusted Fourier transform (AAFT) surrogates)

Generate N realisations of a Gaussian distribution $\{g_t\}$ and reorder $\{g_t\}$ so that it has the same rank distribution as $\{x_t\}$ – i.e. if x_i is the m_i -th largest among all the observations $\{x_t\}$ then g_i should be the m_i -th largest of the $\{g_t\}$ $\forall i$. Generate an Algorithm 1 surrogate of g_t – \hat{g}_t . Create the surrogate s_t by reordering the original data to have the same rank distribution as \hat{g}_t . I.e. $s_t = x_{\tau(t-1)+1}$ where the permutation τ is chosen so that when \hat{g}_i is the m_i -th largest among all the observations $\{\hat{g}_t\}$ then s_i should be the m_i -th largest of the $\{s_t\}$

Algorithm 2 tests the hypothesis that the observed data is *monotonic static nonlinear transform of linearly filtered noise* This is useful when the data is clearly not Gaussian, but one suspects the underlying cause to be a linear stochastic process. To address this, we rescale the data to be Gaussian, apply Algorithm 1, and then rescale the output to match the original data.

- Algorithm 0 preserves, in the surrogates, the probability distribution of the data.
- Algorithm 1 preserves the linear correlation structure (and Fourier power spectrum) of the data
- Algorithm 2 preserves, approximately, both probability distribution and power spectrum

Testing your data with surrogates

example....

Further reading

Further reading

- B. Thorne, T. Jüngling, M. Small and M. Hodkiewicz. “Parameter Extraction with Reservoir Computing: Nonlinear Time Series Analysis and Application to Industrial Maintenance” Chaos (2021), in press.
- D.C. Corrêa, J.M. Moore, T. Jüngling and M. Small. “Constrained Markov order surrogates” Physica D **406** (2020): 132437. pdf
- X. Peng, M. Small, Y. Zhao, J.M. Moore. “Detecting and Predicting Tipping Points” International Journal of Bifurcations and Chaos **29** (2019): 1930022 pdf
- K. Sakellariou, T. Stemler and M. Small. “Markov modelling via ordinal partitions: A novel paradigm for network-based time series analysis” Physical Review E **100** (2019), 062307 pdf

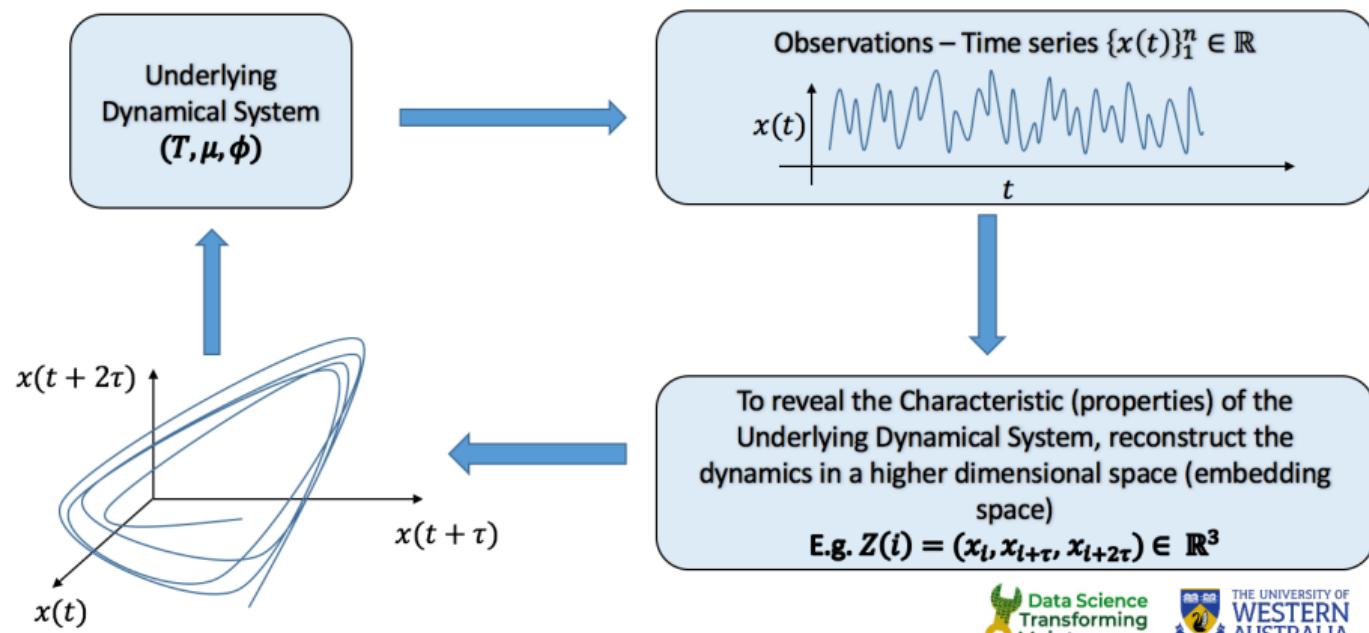
Part II

Pattern Recognition and Change Point Detection Using Recurrence and Dynamical Systems

Embedding review

Embedding Definition

For a time series $\{x(t)\}_1^N$ we define an embedded sequence with embedding dimension m and time-lag τ as $\{Z(k) = (x(k), x(k + \tau), x(k + 2\tau), \dots, x(k + (m - 1)\tau))\}_1^M$.



Embedding Parameters

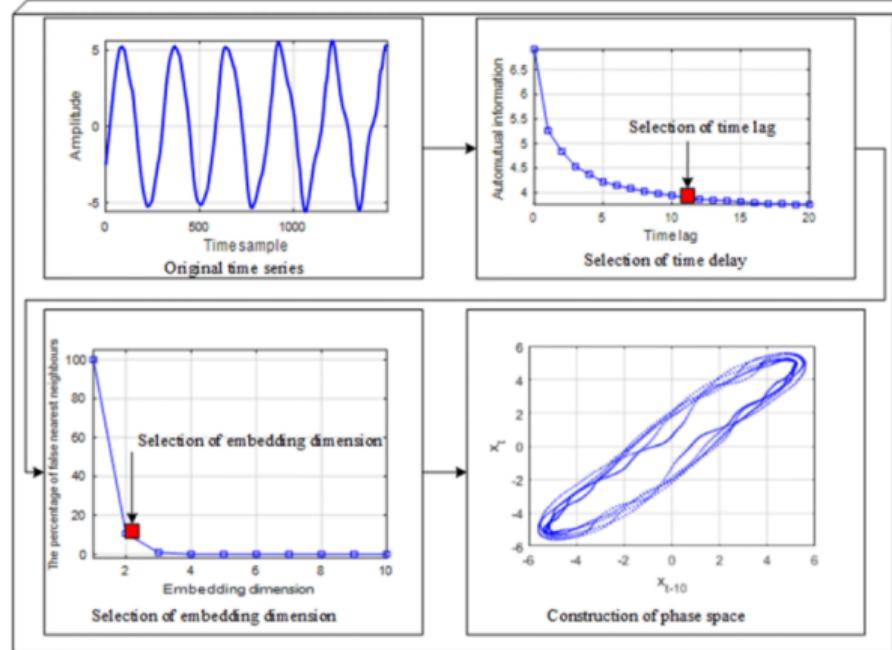
The embedding parameters can be selected by one of the methods⁴:

- For τ :
 - First zero Auto-correlation
 - First minimum Mutual Information
- For m :
 - False Nearest Neighbours (FNN): For each m we count the points that they are close (neighbours) to $x(i)$, but they are not neighbours to $x(i + 1)$

⁴There are other methods addressed by M. Small, *Applied nonlinear time series analysis: applications in physics, physiology and finance*, Vol. 52 (World Scientific, 2005)

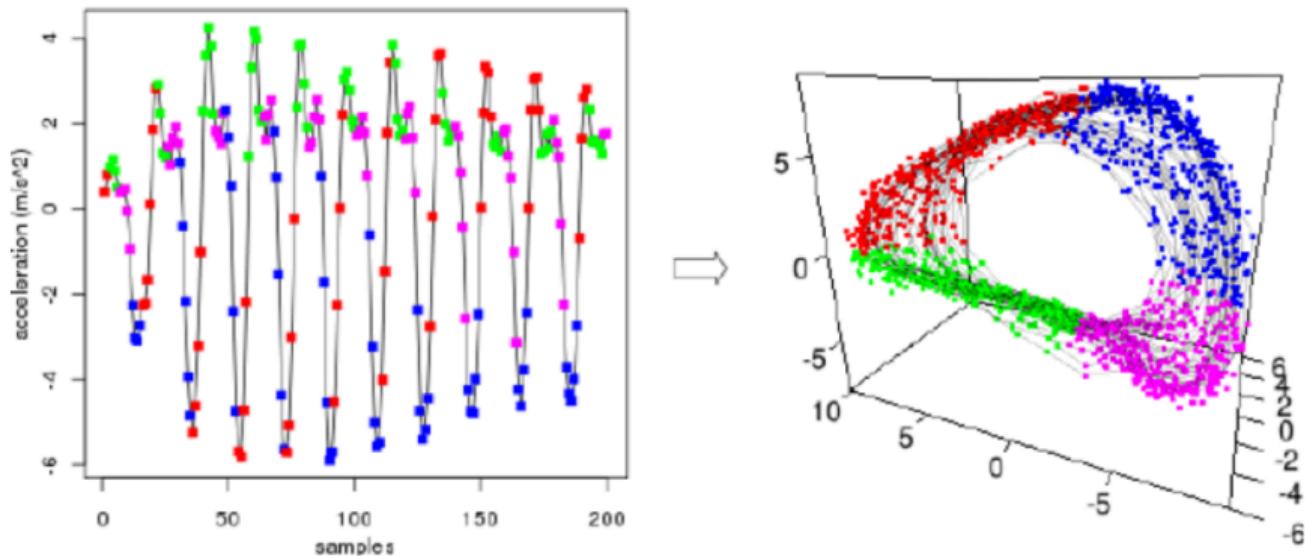
Embedding Parameters

Aydin, I., Karakose, M. & Akin, E. A new method for time series classification using multi-dimensional phase space and a statistical control chart. *Neural Comput & Applic* 32, 7439–7453 (2020). <https://doi.org/10.1007/s00521-019-04270-1>



Embedding Example

Frank, J., Mannor, S. and Precup, D., 2011, September. Activity recognition with mobile phones. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases* (pp. 630-633). Springer, Berlin, Heidelberg.



Recurrence Plots

Why Recurrence Plots?

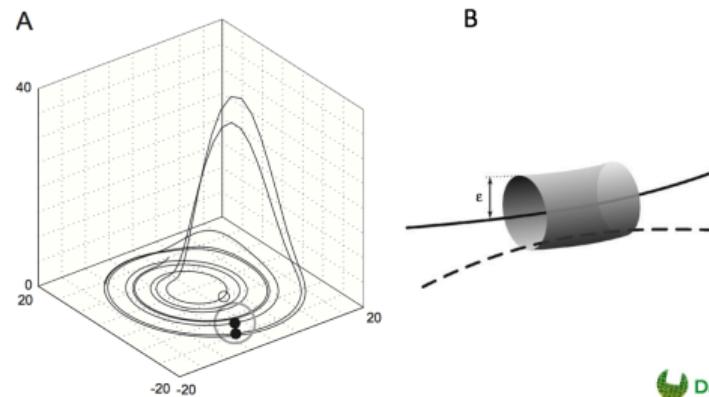
- Recurrences is originally introduced by Henry Poincare in 1890. Later in 1987, Eckmann et al. introduced the method of RPs to visualise the recurrences of dynamical systems.
- Recurrence is a fundamental characteristic of many dynamical systems.
- Recurrence event is defined when two states of the system pass close to each other in different times.
- RPs is a visualisation tool associated with quantitative analysis for time series of nonlinear dynamical systems.
- It provides an alternative and powerful mathematical framework to study time series data and extract features of different systems or patterns.

RPs definition

For a time series (embedded or multivariate) $\mathbf{x}_{i=1}^N \in \mathbb{R}^d$, the Recurrence Plot matrix is defined as follows:

$$RP_{ij} = \begin{cases} 1 & \text{if } \|\mathbf{x}_i - \mathbf{x}_j\| < \varepsilon \\ 0 & \text{otherwise} \end{cases}$$

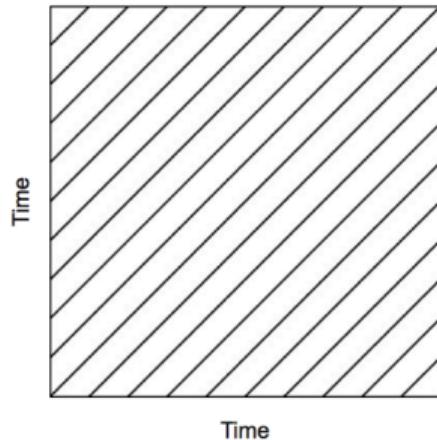
where $\|\cdot\|$ is some norm (distance) metric, e.g. Euclidean distance. ε is the recurrence threshold, could be topological or metric.



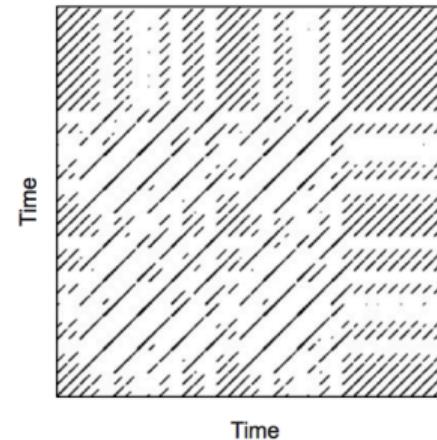
N. Marwan et al. / Physics Reports 438 (2007) 237 – 329

Examples of RPs of some dynamics (patterns)

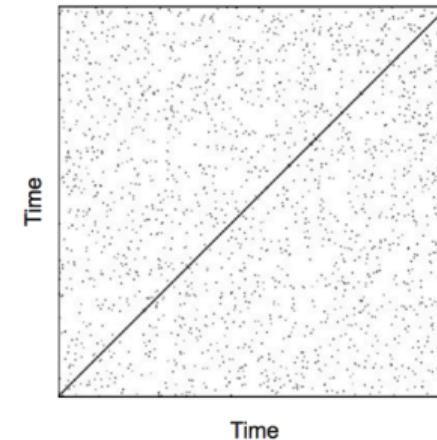
A Periodic system



B Chaotic system

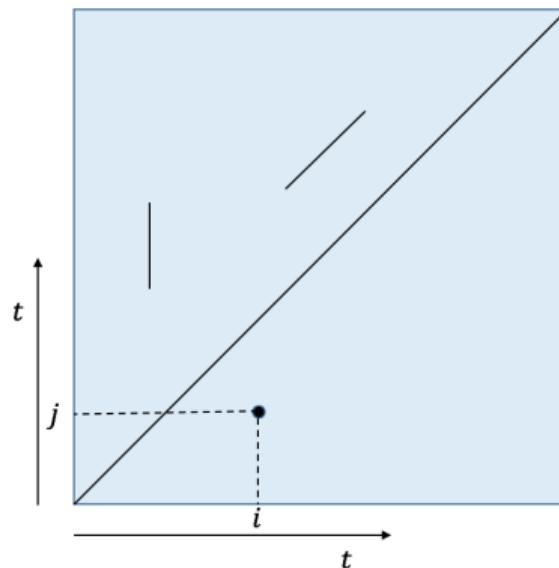


C Uniformly distributed noise

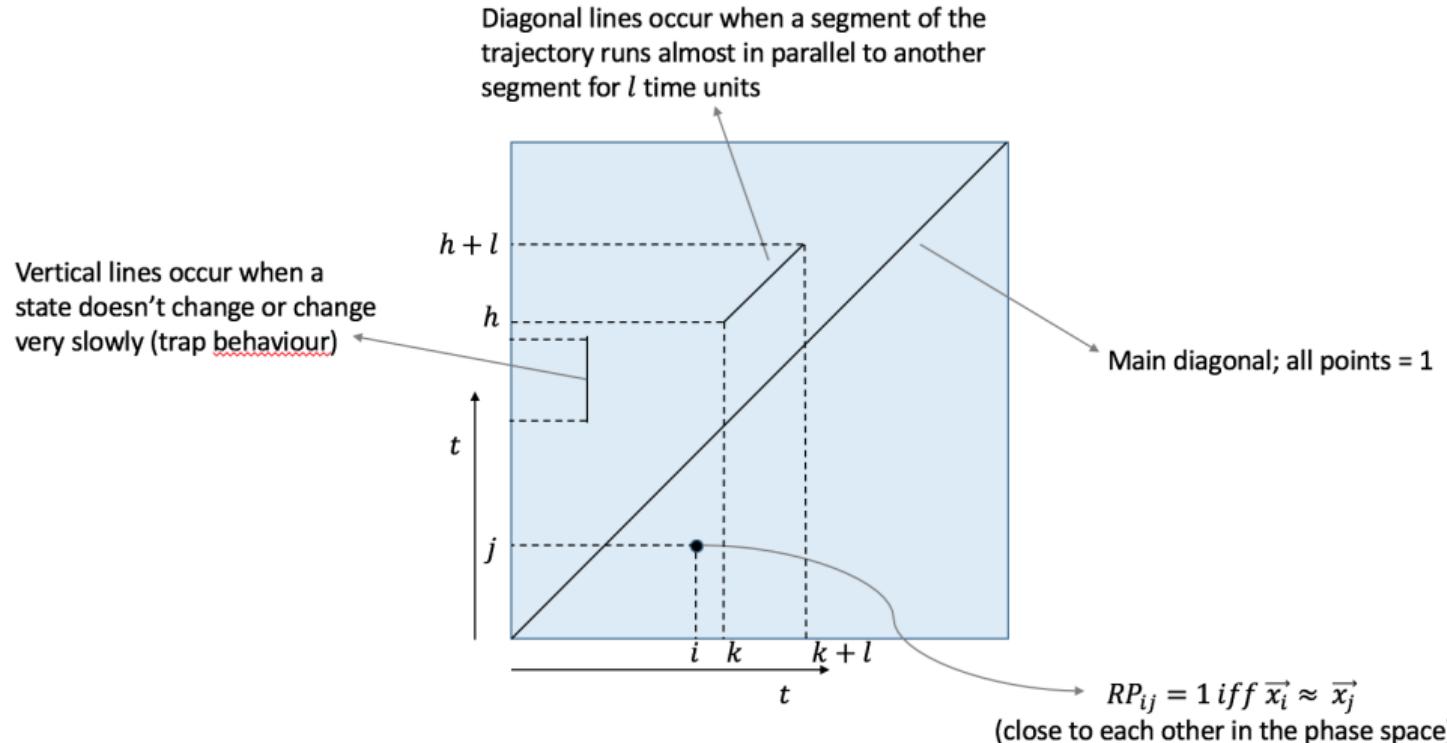


N. Marwan et al. / Physics Reports 438 (2007) 237 – 329

Some structures in RPs

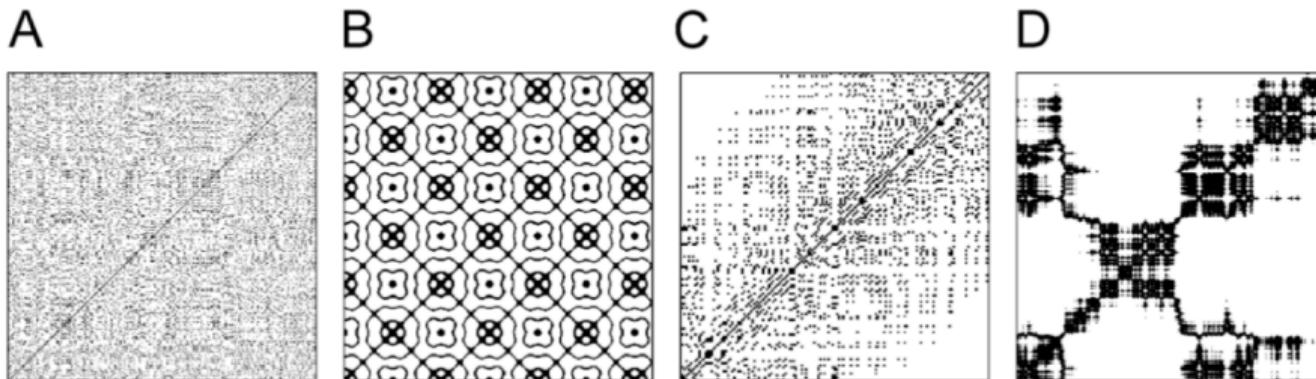


Some structures in RPs



The structure of the RPs reflects important characteristic of the dynamic

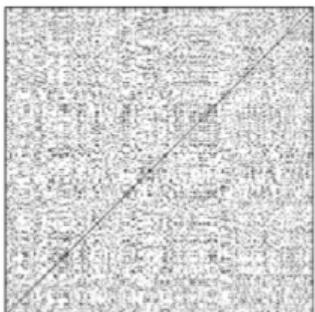
N. Marwan et al. / Physics Reports 438 (2007) 237 – 329



The structure of the RPs reflects important characteristic of the dynamic

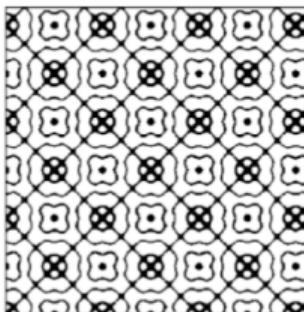
N. Marwan et al. / Physics Reports 438 (2007) 237 – 329

A Homogeneous



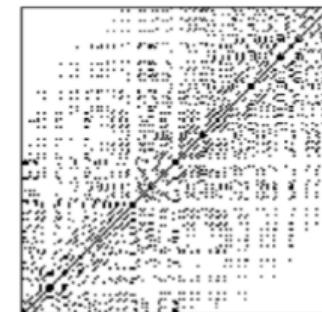
(White noise – stationary
system)

B Periodic



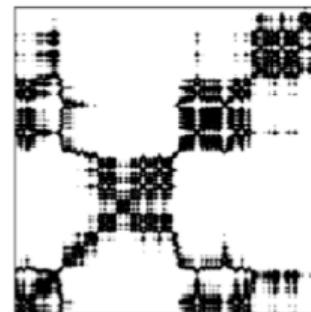
(Super-positioned
harmonic oscillations)

C Drift



(Non-stationary system –
fading to upper left and
lower right corners)

D Abrupt Changes



(White areas – bands)

Typical patterns in RPs

Pattern	meaning
Homogeneity	The system is stationary
Fading to the upper left and lower right corners	The system is non-stationary and contains a trend or a drift
Disruptions (white bands)	The system is non-stationary and abrupt changes may have occurred
Periodic/quasi periodic patterns	The system is periodic, the time distance between the periodic pattern (e.g. diagonal lines) corresponds to the period. Different distances between the patterns reveal quasi-periodic system
Single isolated points	Strong fluctuation in the system
Diagonal segments	The evolution of the system is similar at different epochs. The system could be deterministic. If these diagonal segments beside single isolated points, this means the system could be chaotic.
Vertical or horizontal segments	Some states don't change or change slowly for some time, indication of trap-behaviour (laminar states)

Recurrence Quantification analysis (RQA)

Why RQA?

- RQA provides quantitative measures of the complexity of the time series and the underlying system.
- The measures are based on the density, diagonal lines or vertical lines of the RPs.
- These RQA measures reflect the system's characteristic (properties).
- An alternative and powerful framework to extract features of the system, which can be used for further analysis (e.g. predictive maintenance).

Recurrence Rate (RR): measures the density of recurrence points in the RP. It corresponds with probability that a specific state will recur. (\sim correlation sum).

$$RR = \frac{1}{N^2} \sum_{i,j=1}^N RP_{ij}$$

Where N is the length of the time series.

Determinism (DET): is the percentage of the recurrence points which form diagonal lines of minimal length l_{min} in the RP.

$$DET = \frac{\sum_{l=l_{min}}^N lP(l)}{\sum_{l=1}^N lP(l)}$$

where $P(l)$ is the frequency distribution of length l of the diagonal lines (i.e. it counts how many instances have length l).

DET is related with the predictability of the dynamical system, for example:

- Random system of white noise \Rightarrow RP has almost only single dots and very few diagonal lines \Rightarrow smaller DET .
- Deterministic process \Rightarrow RP has very few single dots but many long diagonal lines \Rightarrow larger DET .

Average diagonal line length (L): is the average length of the diagonal lines.

$$L = \frac{\sum_{l=l_{min}}^N l P(l)}{\sum_{l=l_{min}}^N P(l)}$$

It is related with the predictability time of the dynamical system (i.e. the average time that two segments of the system's trajectory are close to each other).

Divergence (DIV): is the inverse of the maximal diagonal line length L_{max} .

$$DIV = \frac{1}{L_{max}}$$

where $L_{max} = \max(\{l_i\}_{i=1}^{N_l})$, where N_l is the total number of diagonal lines in the RP.

DIV is related with the positive Lyapunov exponent of the dynamical system (i.e. faster trajectory segments diverge \Rightarrow shorter are the diagonal lines \Rightarrow larger DIV).

Entropy (ENTR): is the Shannon entropy of the probability of the diagonal line lengths $p(l)$.

$$ENTR = - \sum_{l=l_{min}}^N \rho(l) \ln(\rho(l))$$

where $\rho(l) = \frac{P(l)}{\sum_{l=l_{min}}^N P(l)}$ (the probability that a diagonal line has length l).

ENTR reflects the complexity of the RP in respect of the diagonal lines.

Ratio: is the ratio between *DET* and *RR*.

$$RATIO = N^2 \frac{\sum_{l=l_{min}}^N lP(l)}{(\sum_{l=1}^N lP(l))^2}$$

It can be used to uncover transitions in the dynamics.

Laminarity (LAM): is the percentage of the recurrence points which form vertical lines of minimal length v_{min} in the RP.

$$LAM = \frac{\sum_{v=v_{min}}^N vP(v)}{\sum_{v=1}^N vP(v)}$$

where $P(v)$ is the frequency distribution of length v of the diagonal lines (i.e. it counts how many vertical lines have length v).

LAM is related with the amount of laminar states in the system (i.e. the states which are trapped for some time).

Trapping Time (TT): is the average length of the vertical lines.

$$TT = \frac{\sum_{v=v_{min}}^N vP(v)}{\sum_{v=v_{min}}^N P(v)}$$

TT is related with the laminarity time of the dynamical system (i.e. how long the system remains in a specific state).

Maximal vertical line length (v_{max}): is the longest time in which a state is trapped.

$$v_{max} = \max(\{v_l\}_{l=1}^{N_v})$$

where N_v is the total number of vertical lines in the RP.

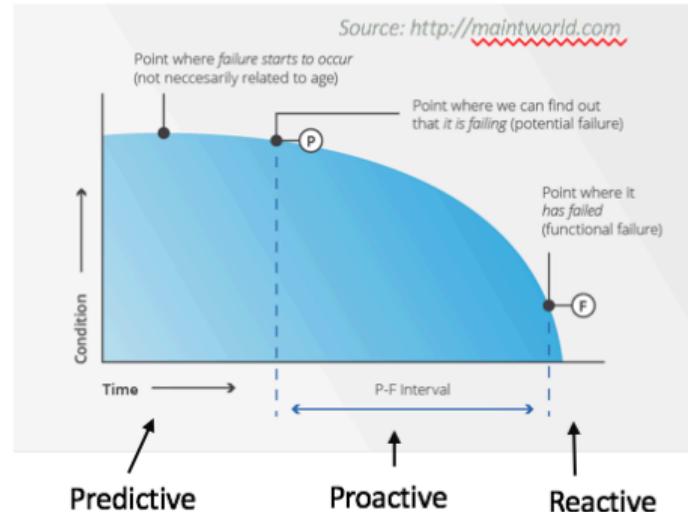
RPs to detect change points

Motivation



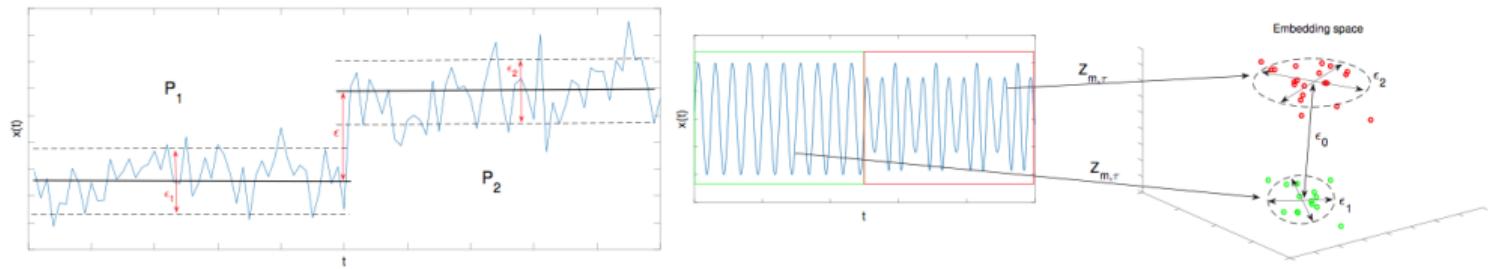
Source: <http://lase.de>

Source: <http://tenova.com>



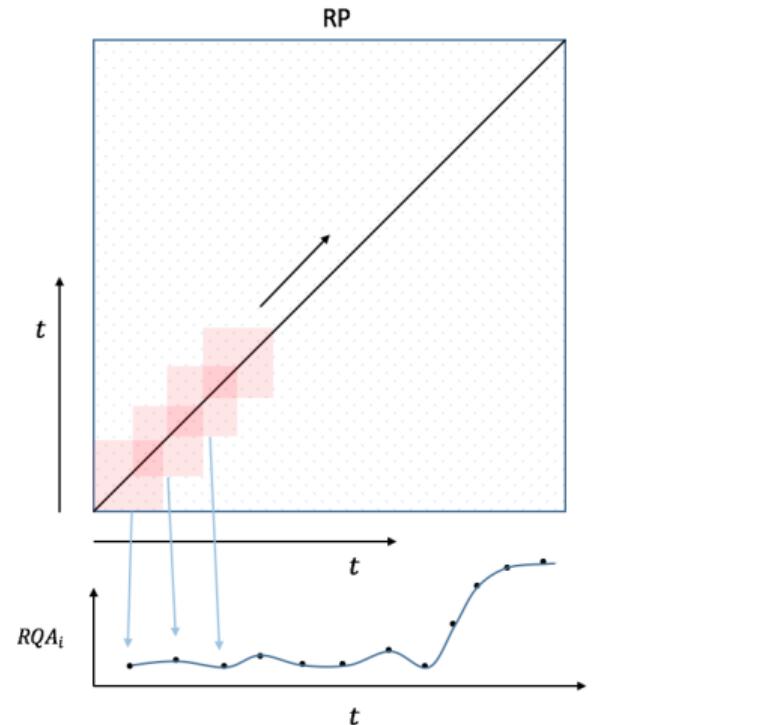
Change points (tipping points)

- Tipping points are the critical points at which an important change in the system occurs or is derived (e.g. a bifurcation value of a system parameter).
- Locating or detecting these tipping points is of high importance to many fields – including predictive maintenance.
- Two types of transition:
 - State–transition
 - Dynamic–transition



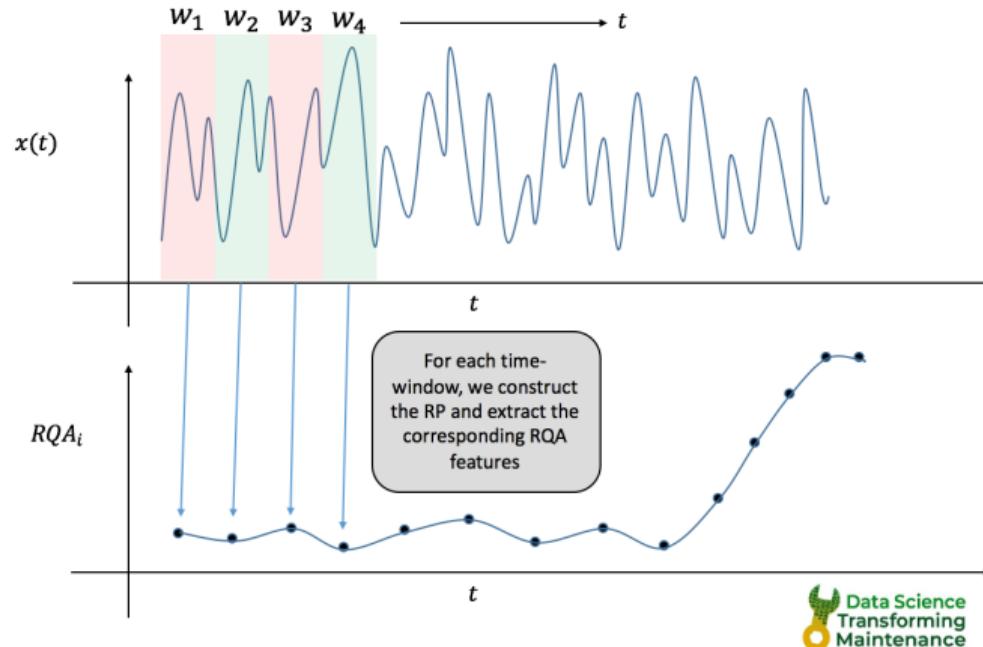
Time-dependent RQA

Approach 1 – sliding window: Instead of computing the RQA measures of the entire RP, RQA measures can be computed in small windows sliding over the RP along the main diagonal \Rightarrow This provides time-dependent RQA measures.



Time-dependent RQA

Approach 2 – windowing the time series: Alternatively, in this approach, the time series is segmented into sequential windows. For each window, a RP is created and associated RQA measures are computed \Rightarrow This provides time-dependent RQA measures.



Quadrant Scan

Definition: From the recurrence plot matrix RP , we construct a sequence $QS(k)$ by counting the ratio of the density of points of those that are in the quadrants with $i, j < k$ or $i, j > k$ versus the density of whole points in all quadrants of RP .

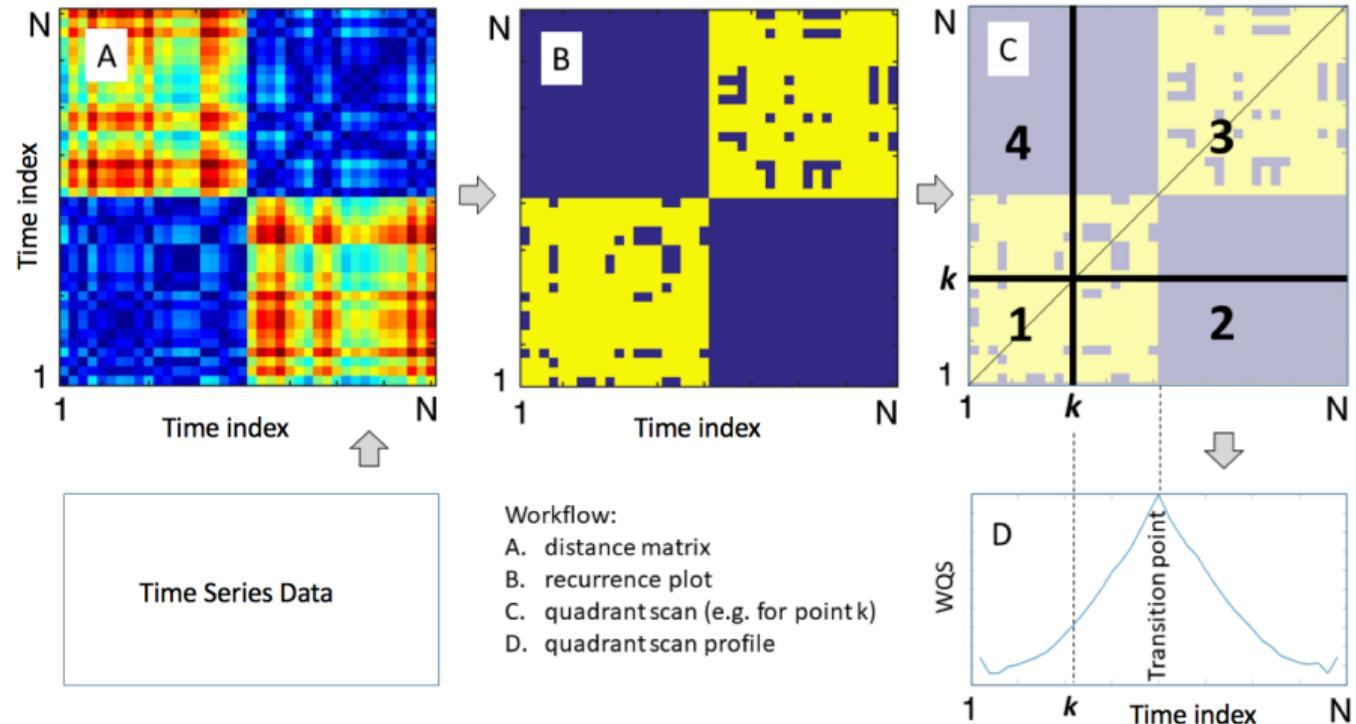
$$QS(k) = \frac{D_{1,3}}{D_{1,3} + D_{2,4}},$$

where $D_{1,3}$ is the density of the points in quadrants 1 and 3 while $D_{2,4}$ is the density in quadrants 2 and 4. They are defined as follows:

$$\begin{aligned} D_{1,3} &= \frac{\sum_{i,j < k} RP_{ij} + \sum_{i,j > k} RP_{ij}}{(k-1)^2 + (N-k)^2} \\ D_{2,4} &= \frac{\sum_{i < k, j > k} RP_{ij} + \sum_{i > k, j < k} RP_{ij}}{(k-1) \times (N-k) \times 2} \end{aligned}$$

Maxima (peaks) of $QS(k)$ correspond to transitions in the system. This is proved in the following theorem. The value of $QS(k)$ is between 0 and 1.

Quadrant Scan

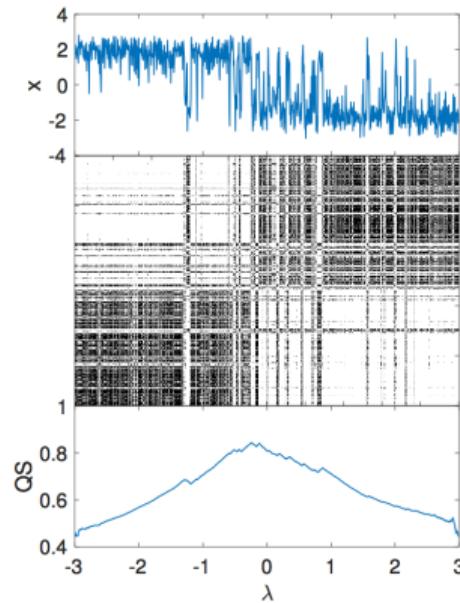
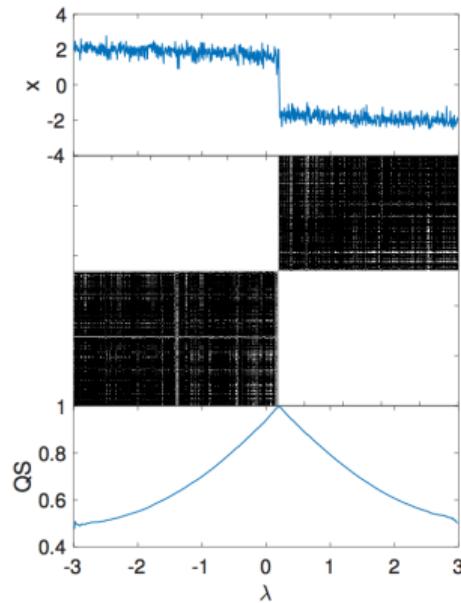
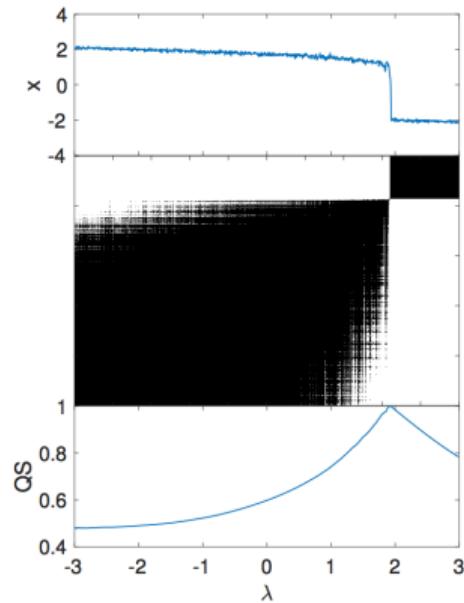


Workflow:

- A. distance matrix
- B. recurrence plot
- C. quadrantscan (e.g. for point k)
- D. quadrantscan profile

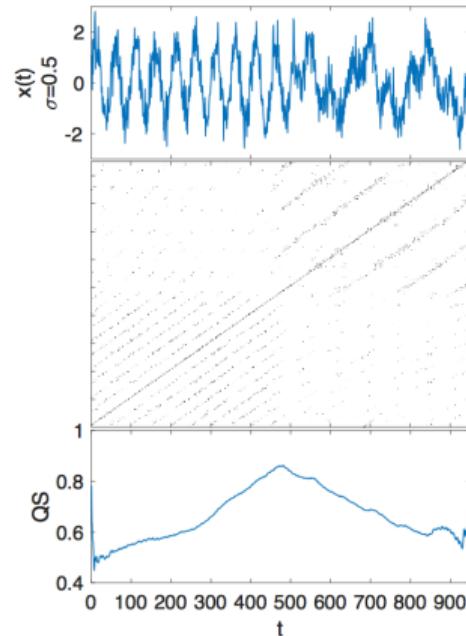
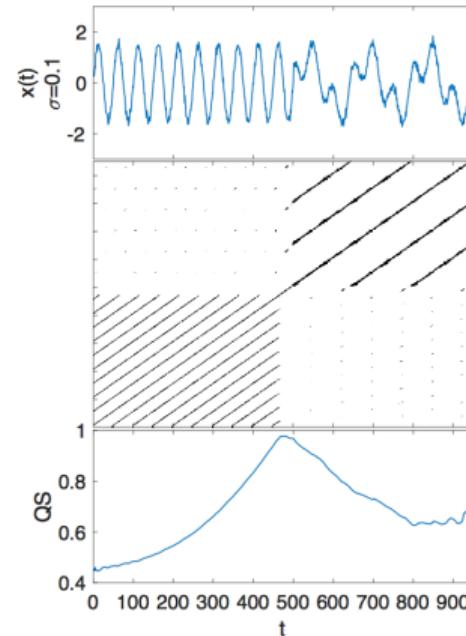
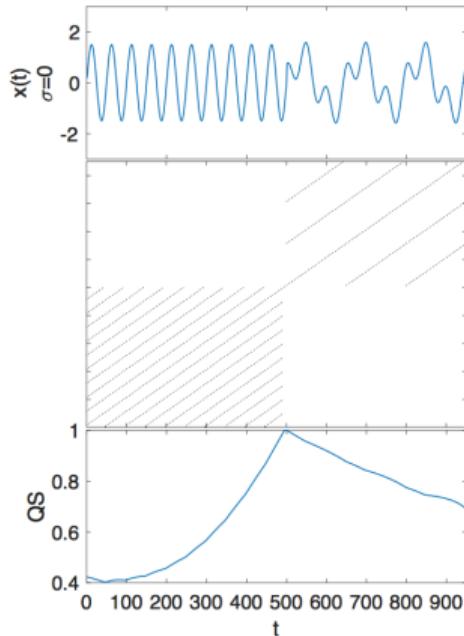
Example of detecting state-transition (type 1)

Noisy Stochastic System:

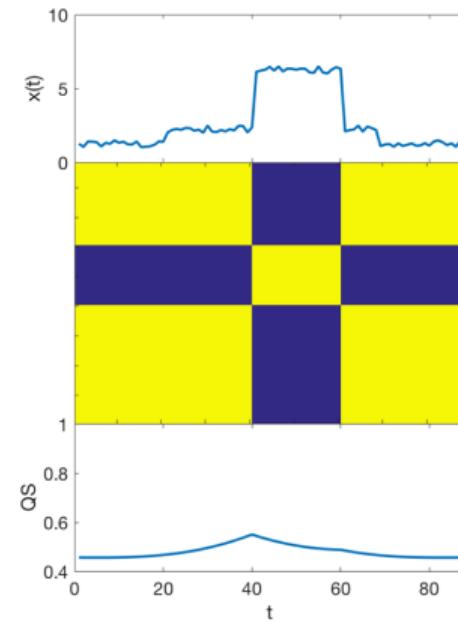
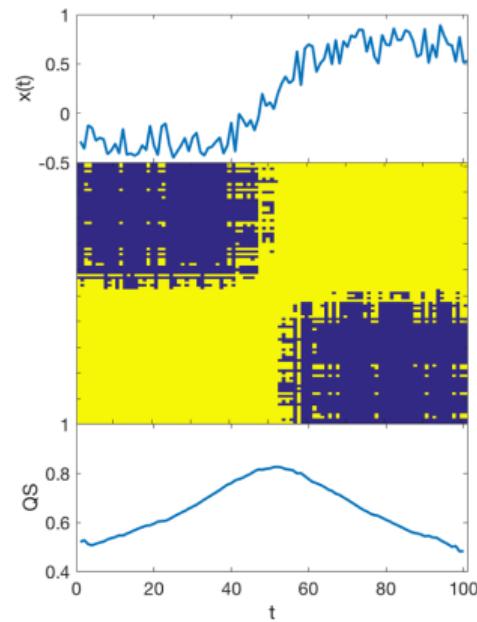
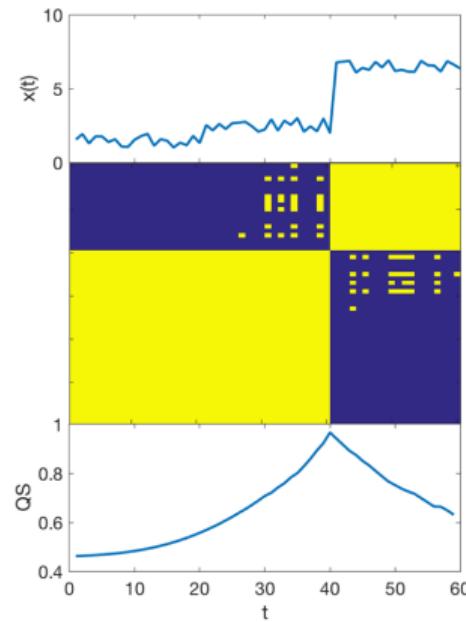


Example of detecting dynamic-transition (type 2)

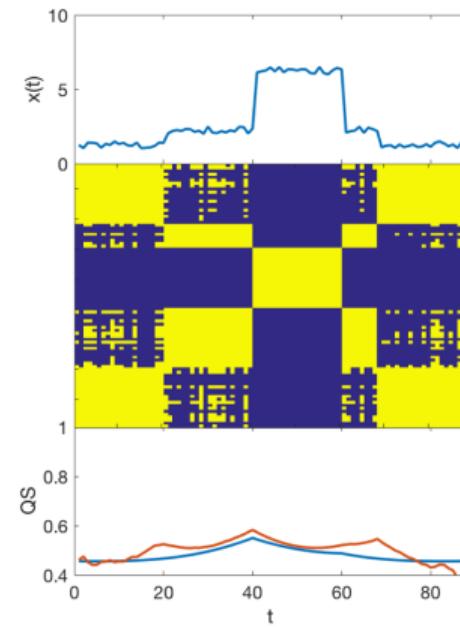
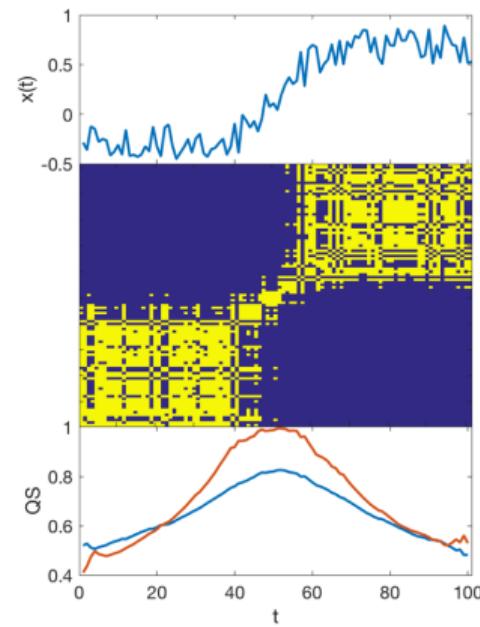
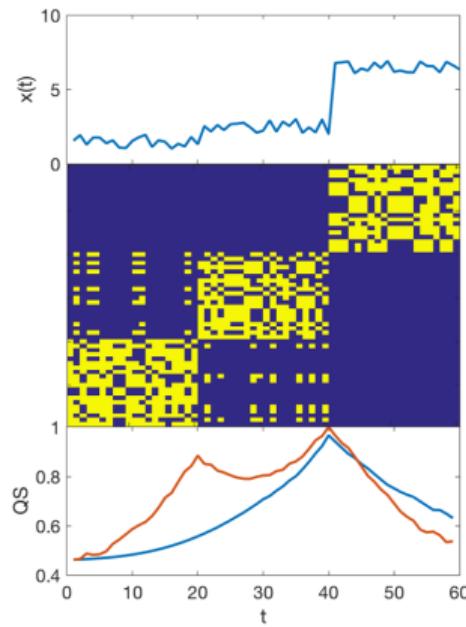
Periodic Stochastic System:



Different scenarios

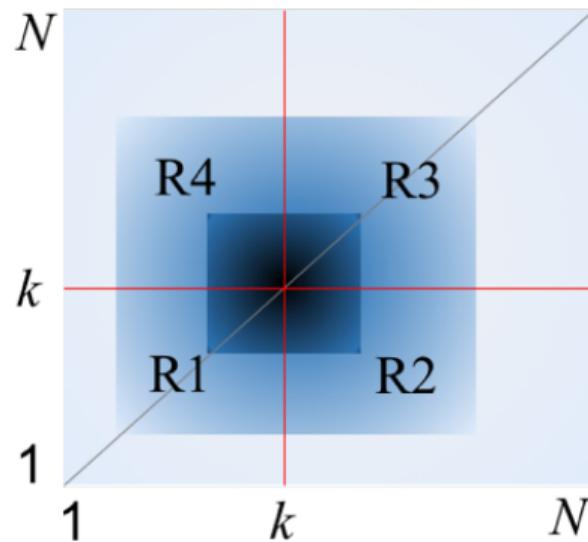


Different scenarios

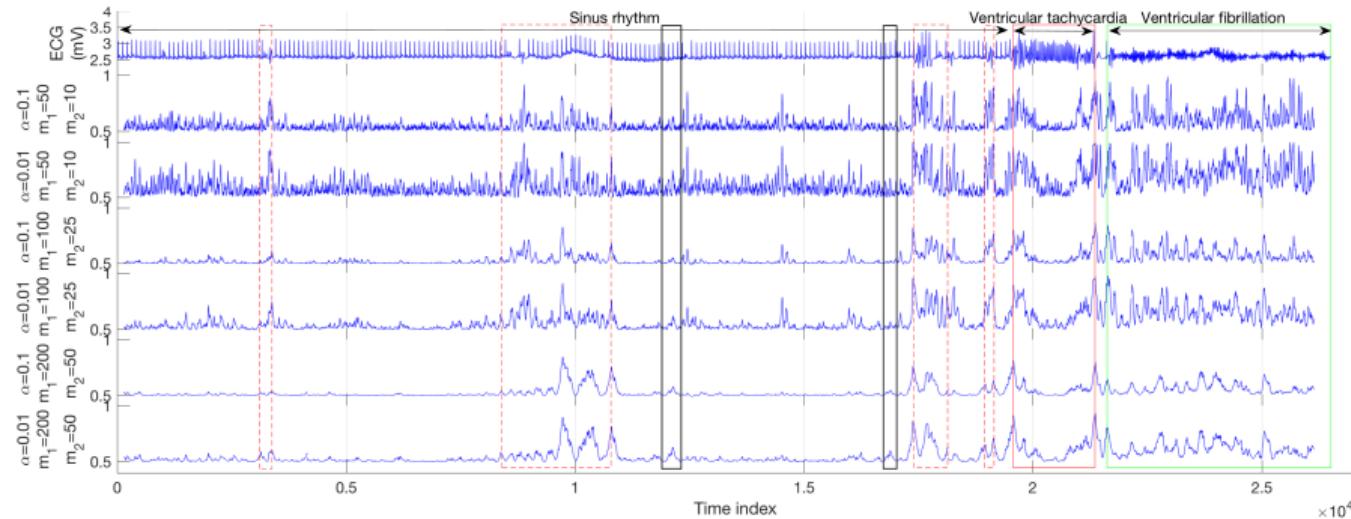


Weighted Quadrant Scan

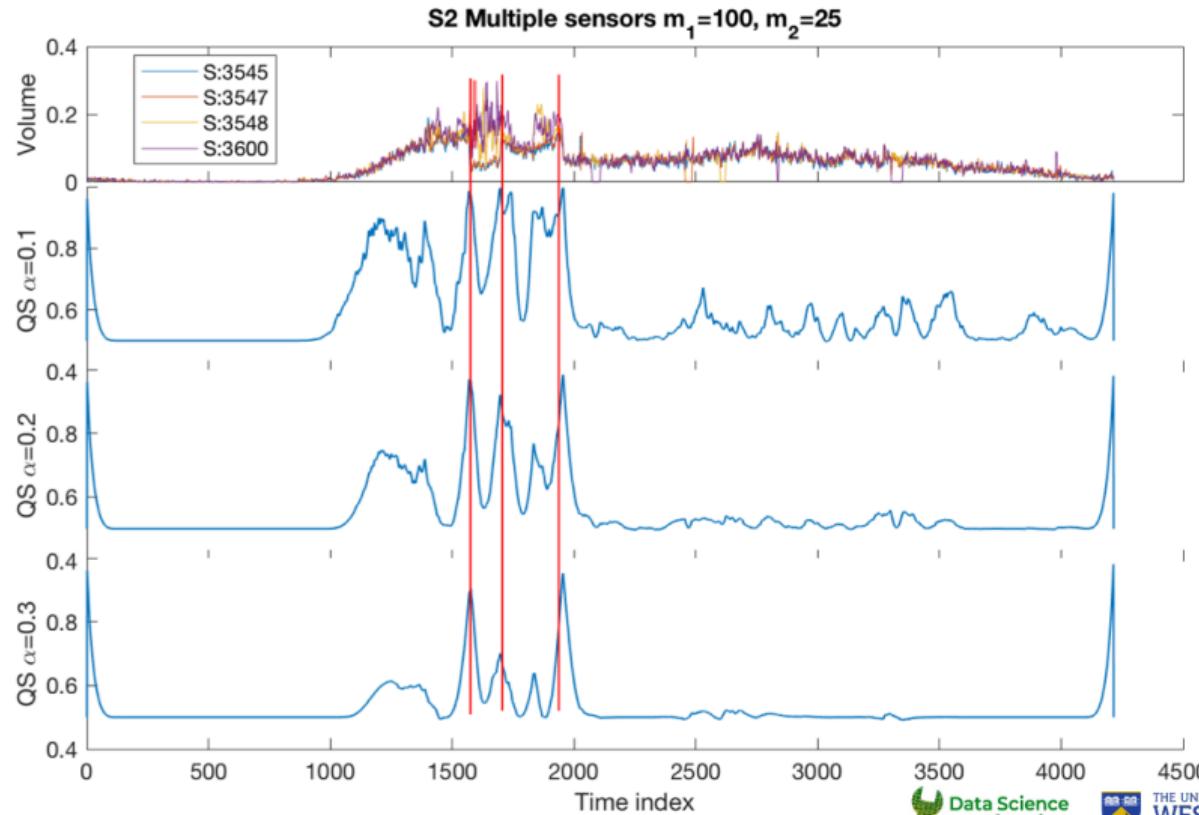
- To detect local transitions, we assign higher weightings to closer points.
- At time k , the dark region represents higher weighted points and the light region for lower weighted points.



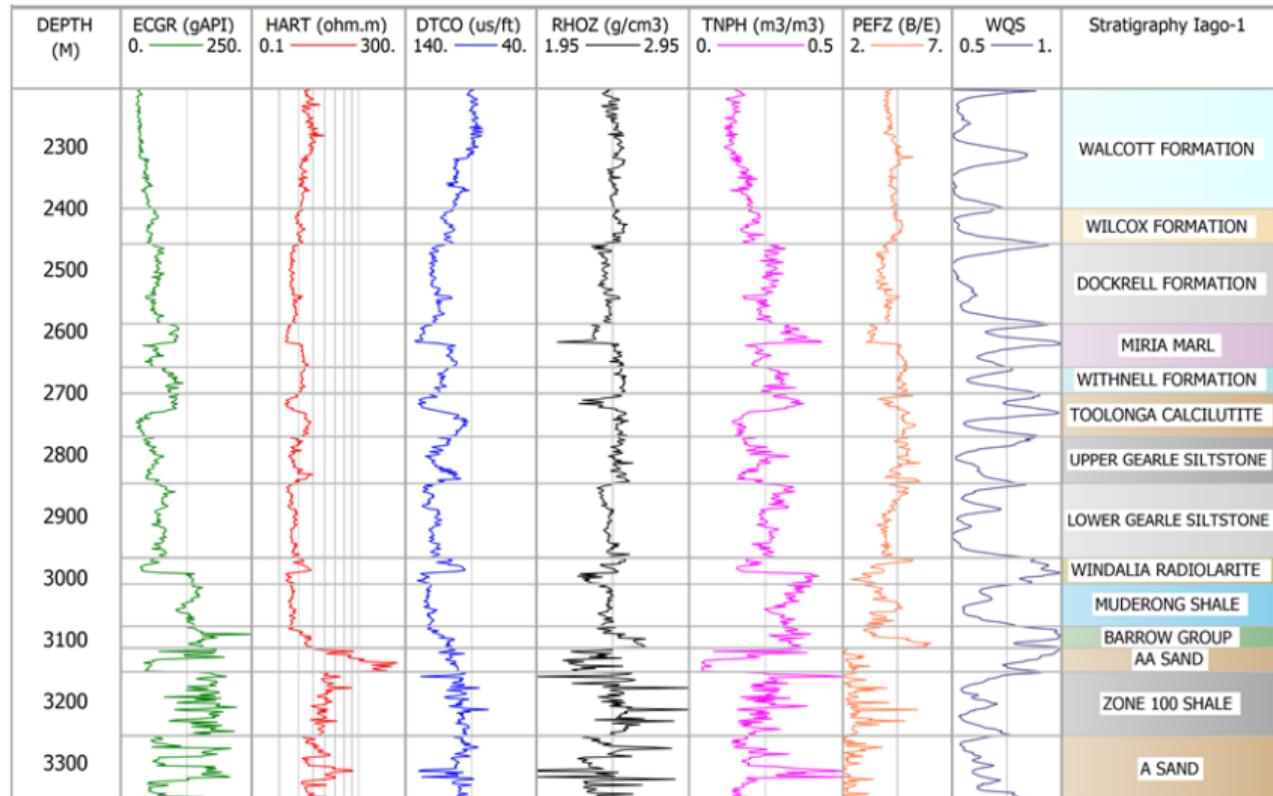
Real-world application 1: ECG data



Real-world application 2: Traffic data



Real-world application 3: Geological data



Part III

Permutation and ordinal representation of data

Recurrence Networks review

Opening of Section, not within a subsection

Ordinal Partition Networks

Section

Permutation Entropy

Section

Application to data

Section

Thank you!

Contact:

Prof. Michael Small

michael.small@uwa.edu.au

Dr. Ayham Zaitouny

ayham.zaitouny@uwa.edu.au

Dr. Débora Corrêa

debora.correa@uwa.edu.au

Acknowledgement

This work is supported by the Australian Research Council through the Industrial Transformation Training Centre for Transforming Maintenance through Data Science (grant number IC180100030) funded by the Australian Government and our industry partners.